

## **1. Source Code For All Programs**

- Pipeline.c
- Plot\_performance.py
- plot\_2.py

## **2. Workload Explanation**

The **Producer-Consumer** workload simulates a **multithreaded pipeline** where multiple producer-consumer pairs interact using **shared buffers**. Each thread pair operates independently to minimize contention.

### **1. Workload Behavior**

- **Producers** generate random data and insert it into a buffer.
- **Consumers** retrieve and process data from the buffer.
- **Synchronization** is managed using **mutex locks** and **condition variables** to ensure thread-safe access.

Each thread performs **100 operations** (producing or consuming items). The buffer size is **100**, preventing buffer overflow or underflow.

### **2. Characteristics of the Workload**

- **Multi-threading & Parallelism**: Each producer-consumer pair runs independently, increasing parallelism.
- **Synchronization Overhead**: Mutex locks and condition variables introduce delays due to waiting threads.
- **Memory Access Patterns**: Frequent buffer read/write operations impact cache efficiency.
- **Scalability Challenge**: Performance may degrade with excessive threads due to increased context switching.

### **3. Objective of Performance Analysis**

- Measure **execution time** across different thread counts.
- Analyze **throughput (items/sec)** to evaluate efficiency.
- Investigate **context switching, cache performance, and CPU cycles** using **perf**.

## **3. Perf Analysis**

The **Producer-Consumer** workload was analyzed using the **perf** tool to measure key system performance metrics. The evaluation was conducted across varying thread counts to study how synchronization and parallelism impact execution.

- **Key performance metrics observed–**

1. **Execution Time**

- Measured total runtime for different thread counts.
- Decreases initially with more threads but stabilizes due to synchronization overhead.

2. **Throughput (items/sec)**

- Number of processed items per second.
- Peaks at an optimal thread count before declining due to contention.

3. **Context Switches**

- Indicates scheduling overhead.
- Increases with higher thread counts due to mutex-based synchronization.

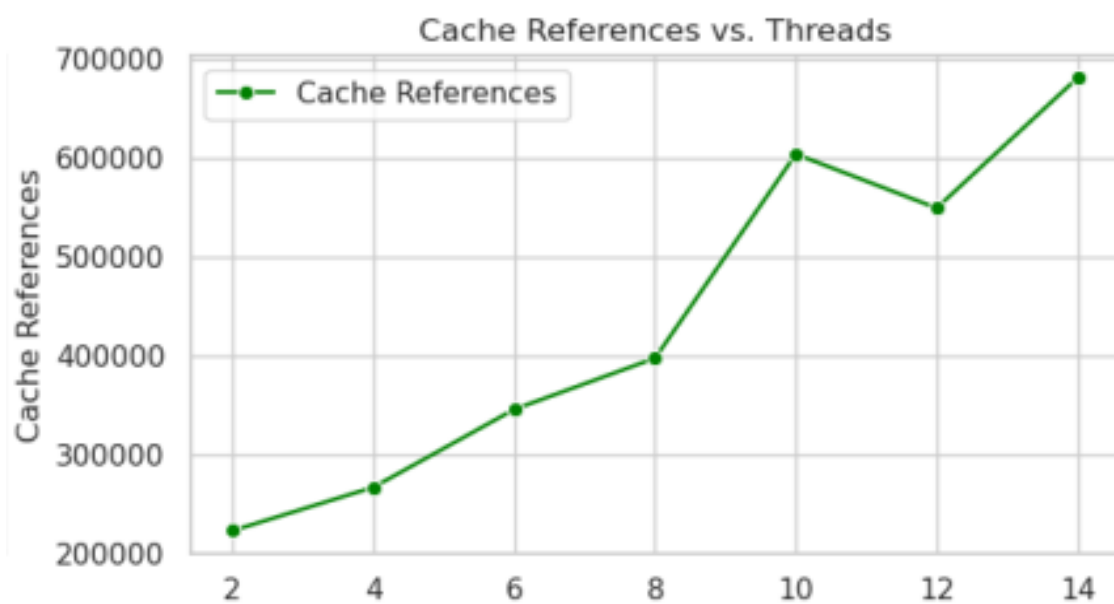
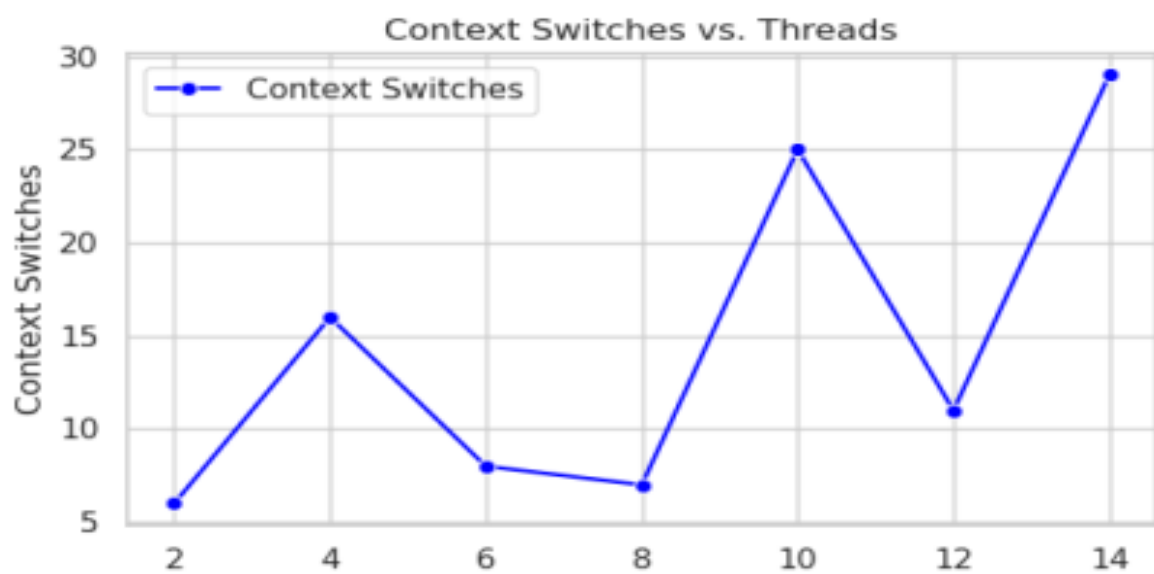
4. **Cache References & Misses**

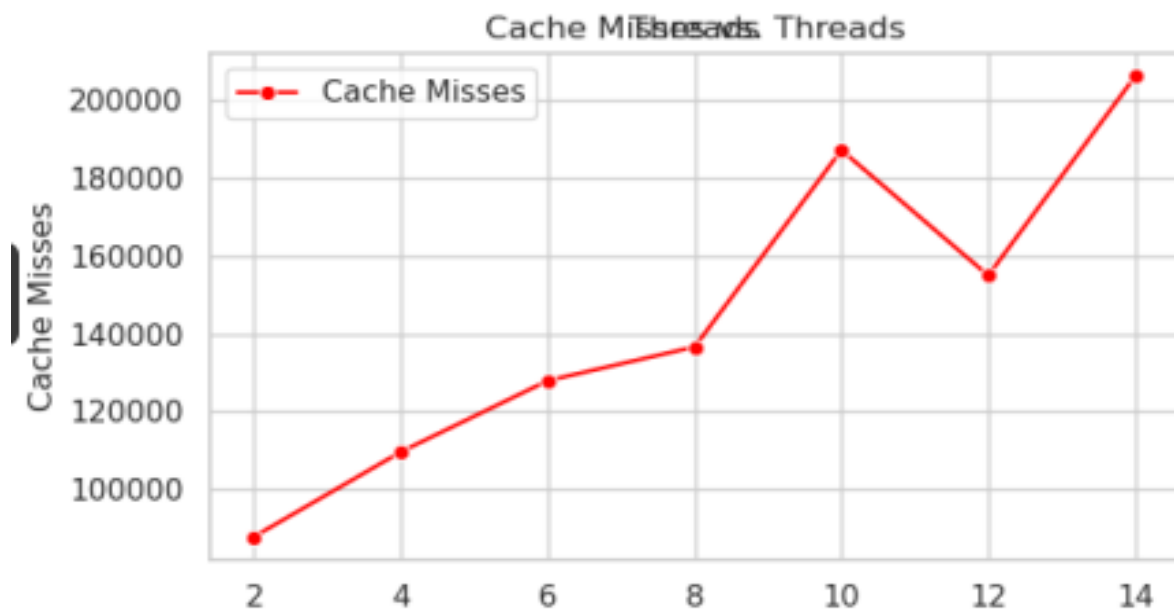
- High cache misses indicate inefficient memory access patterns.
- Shared buffer accesses lead to frequent cache invalidations.

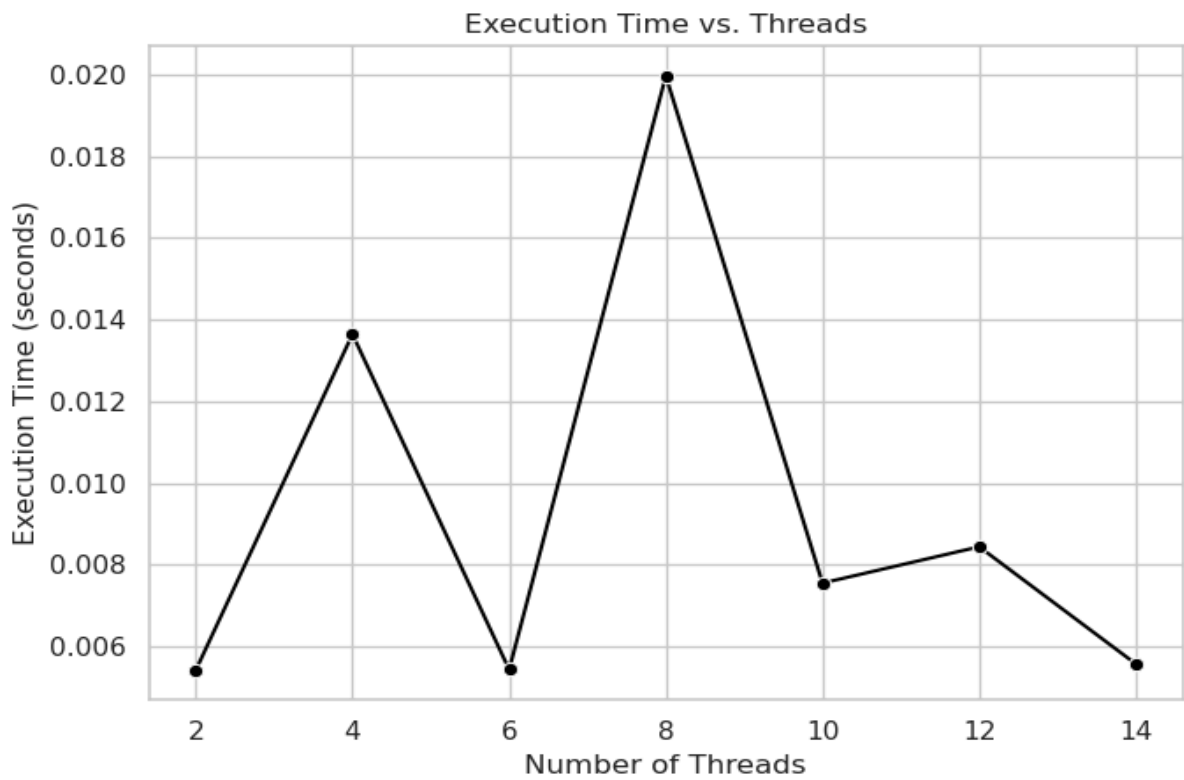
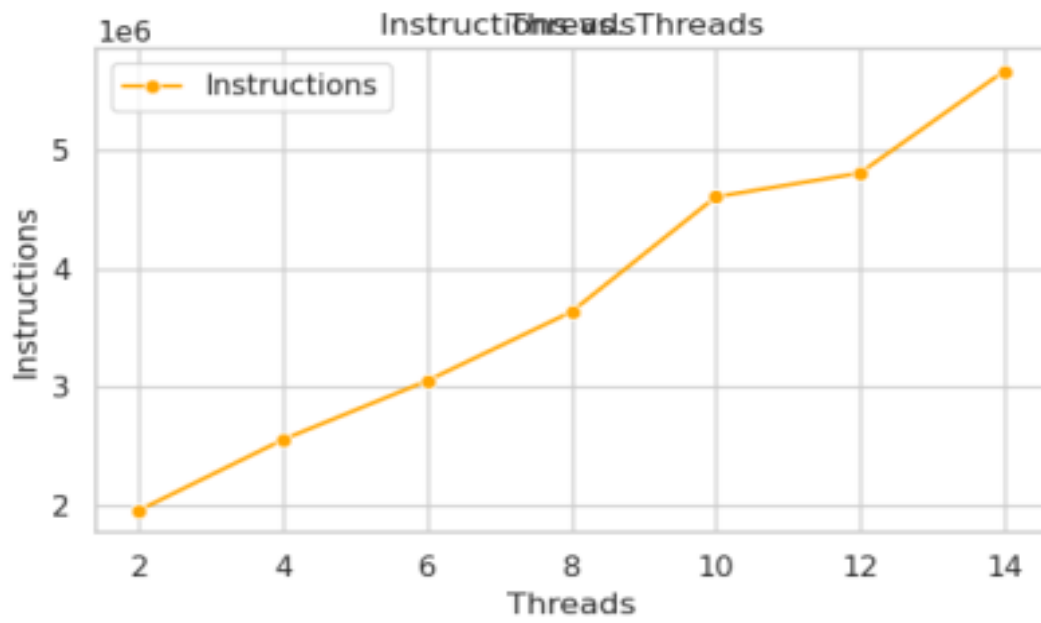
5. **CPU Cycles & Instructions**

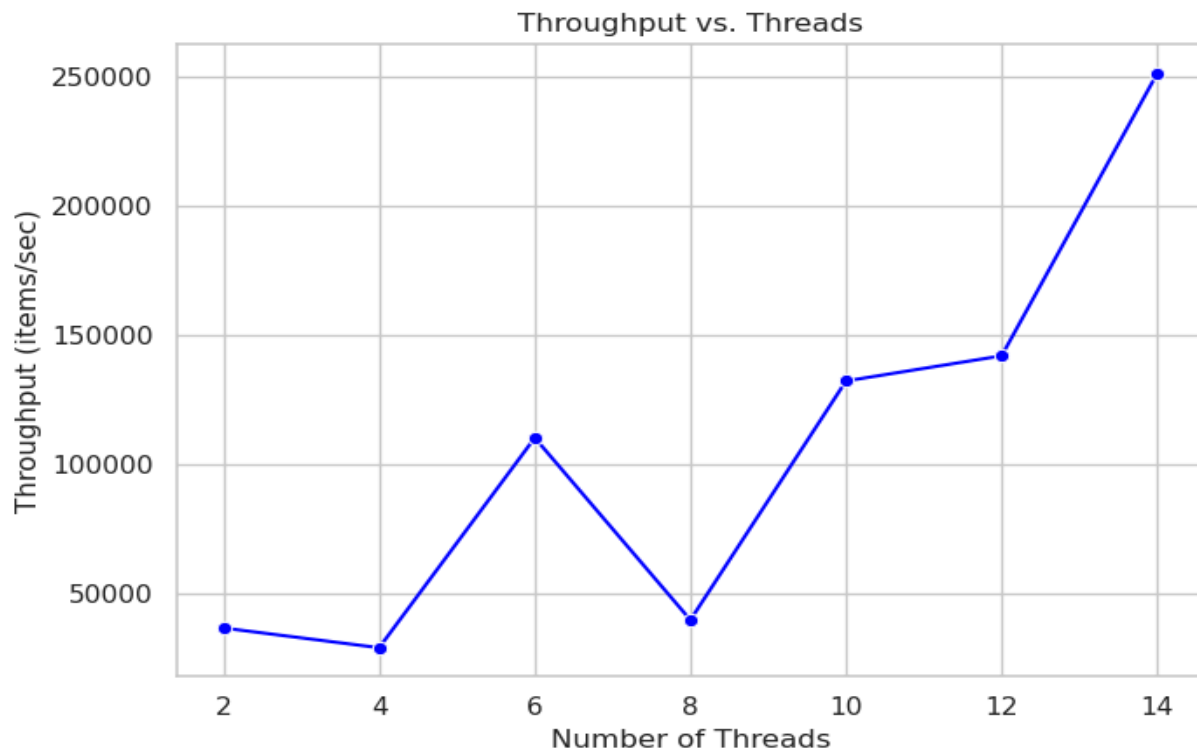
- Helps analyze CPU-bound performance.
- Low instructions per cycle (IPC) suggest waiting on locks rather than execution.

## 4. **Graphs-**









## 5. Screenshots of perf tool outputs for each program.

```
= perf_2.txt
1 Execution Time: 0.005414 seconds
2 Throughput: 36941.26 items/sec
3
4 Performance counter stats for './pipeline 2':
5
6           |   |   |   |   |   | 6      context-switches
7             222,782    cache-references
8              87,906    cache-misses          #    39.46% of all cache refs
9            8,013,472    cycles
10           1,954,127    instructions         #     0.24 insn per cycle
11
12        0.011720948 seconds time elapsed
13
14        0.001115000 seconds user
15        0.003347000 seconds sys
16
17
```



```

perf_10.txt
1 Execution Time: 0.007547 seconds
2 Throughput: 132502.98 items/sec
3
4 Performance counter stats for './pipeline 10':
5
6      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
7      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
8      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
9      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
10     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
11     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
12     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
13     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
14     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
15     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
16     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
17

```

Event	Value	Unit	Notes
context-switches	25		
cache-references	603,307		
cache-misses	187,160		# 31.02% of all cache refs
cycles	18,952,759		
instructions	4,604,097		# 0.24 insn per cycle
seconds time elapsed	0.009369498		
seconds user	0.001422000		
seconds sys	0.005689000		

```

perf_12.txt
1 Execution Time: 0.008437 seconds
2 Throughput: 142230.65 items/sec
3
4 Performance counter stats for './pipeline 12':
5
6      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
7      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
8      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
9      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
10     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
11     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
12     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
13     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
14     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
15     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
16     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
17

```

Event	Value	Unit	Notes
context-switches	11		
cache-references	548,790		
cache-misses	154,947		# 28.23% of all cache refs
cycles	11,280,430		
instructions	4,804,881		# 0.43 insn per cycle
seconds time elapsed	0.007223084		
seconds user	0.001005000		
seconds sys	0.003017000		

```

perf_14.txt
1 Execution Time: 0.005566 seconds
2 Throughput: 251527.13 items/sec
3
4 Performance counter stats for './pipeline 14':
5
6      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
7      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
8      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
9      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
10     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
11     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
12     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
13     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
14     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
15     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
16     |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |      |
17

```

Event	Value	Unit	Notes
context-switches	29		
cache-references	680,232		
cache-misses	206,057		# 30.29% of all cache refs
cycles	17,472,720		
instructions	5,667,387		# 0.32 insn per cycle
seconds time elapsed	0.015733204		
seconds user	0.000000000		
seconds sys	0.007264000		



## **6. Discussion of results, including bottlenecks and optimization opportunities.**

The performance analysis of the Producer-Consumer workload revealed several key insights regarding multithreading efficiency, synchronization overhead, and resource utilization.

- **Bottlenecks–**

1. **Synchronization Overhead**

- The use of mutexes and condition variables leads to contention when multiple threads attempt to access the shared buffer.
- As thread count increases, the time spent waiting for locks rises, reducing the benefits of parallelism.

2. **Context Switching Overhead**

- Higher thread counts increase the number of context switches, adding scheduling overhead.
- Excessive context switching leads to wasted CPU cycles, negatively impacting performance.

3. **Cache Performance Issues**

- Shared buffer access leads to cache invalidations and higher cache misses, especially at larger thread counts.
- Increased memory contention results in more frequent data movement between caches and RAM, slowing execution.

4. **Throughput Saturation**

- Throughput increases initially but plateaus beyond an optimal number of threads.
- At very high thread counts (e.g., 50+ threads), throughput declines due to lock contention and scheduling overhead.

## **Optimization Opportunities**

1. **Reducing Lock Contention**

- Use **lock-free data structures** (e.g., atomic operations, concurrent queues) to minimize waiting time.
- Implement **fine-grained locking** instead of a single lock per buffer to allow more parallel access.

## 2. **Batch Processing**

- Instead of **locking per item**, producers could insert **multiple items at once**, reducing synchronization frequency.
- Consumers can process batches of items before signaling producers.

## 3. **Improving Cache Efficiency**

- Use **thread-local buffers** before transferring data to the shared buffer to reduce cache invalidations.
- Optimize memory layout to improve **spatial locality**, reducing cache misses.

## 4. **Reducing Context Switches**

- Use **thread pinning** to keep threads on specific CPU cores, reducing cache migration overhead.
- Increase **buffer size** to reduce contention and allow more concurrent processing.